

A Modular Quantum-Stochastic Simulator for Hybrid Variational Learning

Wei Qi Zhang

Undergraduate student, DIRO (Université de Montréal)

July 7, 2026

Abstract

Hybrid quantum–classical learning requires simulation frameworks that are modular, reproducible, and capable of modeling modern stochastic noise. Existing simulators often rely on monolithic abstractions or limited noise configurability, making it difficult to prototype new algorithms or evaluate robustness under perturbations. We introduce a modular quantum–stochastic simulation engine integrating variational circuits, explicit noise processes, hybrid gradient estimation, and reproducible configuration management. Through targeted experiments, we analyze noise sensitivity, optimizer behavior, and hybrid gradient stability. This work provides a lightweight, extensible platform for early-stage QML research and reproducible algorithmic exploration.

1 Introduction

Quantum machine learning (QML) has emerged as a promising research direction, leveraging quantum circuits to tackle learning tasks that challenge classical models. However, practical QML research still relies heavily on simulation, as access to quantum hardware remains limited, noisy, and expensive. Existing simulators provide useful functionality but often impose rigid abstractions, lack flexible noise modeling, or make it difficult to prototype hybrid quantum–classical optimization pipelines with repro-

ducible behavior. To support rapid experimentation, we introduce a modular quantum-stochastic simulator designed for research prototyping. The engine integrates variational circuits, explicit noise processes, and hybrid gradient estimation into a unified framework that is easy to extend, inspect, and integrate with classical machine-learning workflows. Unlike monolithic simulators, our architecture emphasizes transparent data flow, configurable components, and explicit reproducibility, enabling researchers to isolate the effects of noise, optimization choices, and circuit structure.

1.1 Motivation Leading to Implementation

A key motivation for this simulator is the ability to model modern stochastic behavior. Realistic quantum systems exhibit randomness and instability at multiple levels, including measurement noise, gate noise, sampling variance, and decoherence. Our simulator exposes these limitations explicitly, allowing researchers to study how noise interacts with variational optimization and to analyze hybrid gradient behavior under stochastic perturbations. This makes the engine suitable for convergence analysis, robustness testing, and sensitivity assessment in QML pipelines.

Overall, this work contributes a lightweight, extensible, research-oriented simulation engine for hybrid quantum–stochastic learning. It is intended for early-stage experimentation, algorithmic exploration, and reproducible QML studies, particularly in contexts where access to quantum hardware is limited.

2 Related Work

Research in quantum machine learning relies heavily on simulation frameworks. Tools such as PennyLane, Qiskit Aer, and TensorFlow Quantum provide powerful interfaces for circuit construction and automatic differentiation. However, they often abstract away noise processes or restrict how stochasticity can be injected. Variational quantum algorithms (VQAs) have been widely studied, yet reproducible experimentation under noise remains limited due to implicit randomness and opaque execution pipelines.

2.1 Noise Modeling Exploration

Noise modeling has been explored through approaches such as Kraus operators, depolarizing channels, and sampling-based perturbations. However, these methods are typically embedded within larger frameworks, limiting customization and fine-grained control. Hybrid gradient-estimation techniques — including the parameter-shift rule, finite differences, and stochastic estimators — are well established, yet their behavior under noise is not always transparent.

Our simulator differs by providing explicit, modular noise components, hybrid gradient estimation and reproducible configuration management, making it suitable for controlled robustness studies and optimization dynamics.

3 Simulator Architecture

3.1 Design Principles

This simulator is built around three principles: transparency, modularity and reproducibility. Each subsystem handles a single responsibility, allowing researchers to inspect correctness, swap components and extend the original architecture without changing its core functionality.

3.2 Circuit Module

The circuit module provides a minimal, accessible representation of variational quantum circuits. Circuits are defined as parametrized sequences with explicit metadata describing qubit indices, parameter bindings and differentiability. This module supports custom declarative constructions, parameter updates and evaluation under noise.

3.3 Stochastic Noise Module

Noise modules are implemented as pluggable components. Each module defines a sampling rule and transformation applied to circuit outputs. Supported noise types include gate noise, sampling variance, custom perturbations and measurement noise. Explicit noise modeling enables explicit controlled robustness studies.

3.4 Execution Engine

The execution engine coordinates circuit evaluation, noise sampling and measurement. It provides deterministic seed control, batched execution, logging hooks and unified interfaces returning expectation values and measurement samples. All intermediate states are reproducible.

3.5 Hybrid Gradient Module

Gradients are estimated using parameter-shift rules, finite differences or stochastic sampling. The module exposes control over noise sampling during gradient computation, enabling analysis of variance, stability and convergence behavior.

3.6 Optimization Module

Classical optimizers operation on circuit parameters using gradients produced by the simulator. The module supports reproducible update rules, logging hooks and integration with external ML libraries. Custom optimizers can be added by implementing small interfaces.

3.7 Configuration and Reproducibility

Experiments are defined through structured configuration files specifying circuit architecture, noise models, optimizers, seeds and logging behavior. This ensures deterministic implementations and reproducible experimentations.

3.8 Extensibility

New noise models, circuit types or gradient estimators can be added by implementing small interfaces. The simulator is intended for future research use and algorithmic exploration.

4 Stochastic Noise and Hybrid Gradients

4.1 Noise Modeling

Quantum systems exhibit randomness at multiple levels. Our simulator explicitly models noise through sampling-based transformations applied to circuit outputs. Each noise model defines:

- a sampling distribution
- a transformation rule
- a reproducible seed pathway

This allows controlled studies of robustness and sensitivity.

4.2 Gradient Estimation under Noise

Hybrid gradients are computed using parameter-shift rules, finite differences, or stochastic sampling. Noise can be injected during gradient computation, enabling analysis of:

- variance inflation
- optimizer stability
- convergence degradation
- robustness of hybrid methods

4.3 Optimization Pipeline

Gradients feed into classical optimizers, forming a loop. The simulator exposes all intermediate values, enabling detailed inspection of optimization dynamics under noise.

5 Gradient Logic

5.1 Saddle Points

Saddle points arise at critical points where the gradient vanishes while the Hessian exhibits mixed curvature, producing directions of both ascent and

descent. In high-dimensional, nonconvex machine-learning landscapes, such regions destabilize optimization by creating flat or weakly curved neighborhoods in which gradient signals become unreliable, leading to stalled updates, vanishing gradients, or divergence. These phenomena are amplified in quantum machine learning, where stochastic noise and circuit depth further suppress gradient magnitudes. The resulting barren plateaus represent the quantum analogue of saddle-point stagnation, causing optimization to halt as gradients decay toward zero.

5.2 ADAM Optimizer

The simulator employs the ADAM optimizer, which integrates momentum-based smoothing, RMSProp’s adaptive scaling and bias correction mechanisms. Momentum introduces a velocity term that enables the optimizer to traverse regions where gradients vanish, providing sufficient directional persistence to escape saddle-point stagnation and reducing the likelihood of gradient divergence. RMSProp further stabilizes training by normalizing gradients through running estimates of their squared magnitudes, mitigating the impact of stochastic noise. Bias correction ensures that early-step updates remain well-scaled rather than artificially suppressed. Together, these components make ADAM particularly effective in quantum-stochastic settings, where randomness and the prevalence of saddle-like structures demand robust and noise-resilient optimization dynamics.

In summary, classical saddle points and quantum barren plateaus represent analogous forms of gradient stagnation, each causing optimization to halt as curvature flattens and gradient magnitudes collapse. The incorporation of the ADAM optimizer counteracts these effects by providing momentum-driven escape dynamics and adaptive gradient scaling, thereby stabilizing updates under both classical nonconvexity and quantum stochastic noise. Together, these mechanisms enhance the simulator’s robustness, enabling reproducible training behavior and facilitating future extensions without being hindered by noisy or vanishing gradients.

6 Experiments

6.1 Noise Sensitivity

We evaluate a variational circuit under multiple noise regimes and measure expectation variance, gradient variance, and convergence behavior. These metrics allow us to characterize the circuit’s sensitivity to stochastic perturbations and to assess the stability of hybrid optimization under different noise configurations.

6.2 Optimizer Comparison

We compare classical optimizers — including SGD, ADAM, and momentum-based methods — under identical noise conditions. The simulator’s reproducibility guarantees controlled, fair comparisons across optimization strategies.

7 Results

Results indicate that explicit noise modeling significantly affects gradient stability and optimizer behavior. Noise-aware gradient estimation improves convergence in certain regimes, and comparative studies show that optimizer choice interacts strongly with noise characteristics.

8 Limitations and Future Work

The simulator is constrained by classical execution and does not model full quantum-hardware effects such as cross-talk or realistic decoherence. Future work includes GPU acceleration, integration with QPU backends, expanded noise libraries, and support for deeper circuits.

9 Conclusion

We present a modular quantum-stochastic simulator for hybrid variational learning. By exposing noise processes, gradient estimation, and optimizer dynamics explicitly, the system enables reproducible and controlled experimentation. This work provides a lightweight platform for early-stage QML research and algorithmic exploration.

Acknowledgments

I thank the instructors of IFT1065 and MAT1400 for providing foundational tools that supported the early development of this work.